
RAG + LangChain + Groq Gemini Chatbot — Project Report

1. Problem Explanation

AI chatbots generally provide answers based on pre-trained knowledge, lacking the ability to incorporate user-specific documents or datasets dynamically. Also, job seekers often require systems that suggest relevant jobs based on their resumes.

This project builds a hybrid chatbot that:

- Uses Groq's Gemini LLM (`llama3-8b-8192`) for advanced language understanding.
- Answers both general and document-specific queries by combining retrieval from user-uploaded PDFs (RAG approach).
- Includes a job recommendation feature where resumes are matched to a job dataset CSV.
- Provides a unified UI to upload PDFs, resumes, and job datasets, and ask questions or get job suggestions.

2. Use of RAG with LangChain + Groq Gemini

- **Groq Gemini LLM** is accessed via `langchain_groq` with a secure API key from `.env`.
- User PDFs are processed by:
 - Loading with `PyPDFLoader`.
 - Chunking with `RecursiveCharacterTextSplitter`.
 - Embedding chunks with HuggingFace's `all-MiniLM-L6-v2` embedding model (on CPU).

- Indexed into FAISS vector store for fast similarity retrieval.
 - For each user query:
 - Relevant document chunks are retrieved.
 - Conversation history and retrieved context are combined into a prompt.
 - Groq Gemini generates a response based on this enriched context.
 - For job matching:
 - Jobs are loaded from a CSV dataset.
 - Skills are extracted from uploaded resumes.
 - Matching jobs are suggested based on shared skills.
-

3. Folder and Code Walkthrough

```
your_project/
├─ app.py           # Main Streamlit app with UI and flow control
├─ .env             # Stores GROQ_API_KEY securely
├─ requirements.txt # Python dependencies
└─ _modules/       # Modular helper code for maintainability
    ├─ llm.py       # Loads Groq Gemini LLM securely
    ├─ pdf_processing.py # PDF Loading, splitting, embedding, FAISS indexing
    ├─ prompt.py    # Builds prompts with history and context
    ├─ chat.py      # Defines LangChain chain and queries LLM
    └─ job_matching.py # Loads job CSV, extracts skills, matches jobs
```

4. Challenges Faced and Solutions

- **Embedding model loading error:** Solved by forcing HuggingFace embeddings to use CPU (`model_kwargs={"device": "cpu"}`).
 - **UI clutter from multiple uploads:** Unified file uploader with user selection for file type.
 - **Combining conversation history with retrieved context:** Created a prompt builder that appends history and retrieved docs before LLM query.
 - **Skill extraction accuracy:** Extracted skills from resumes by comparing with the job dataset's skill list.
 - **Secure key management:** Used `.env` and `python-dotenv` to safely load API key without hardcoding.
-

5. Summary of Learnings

- Successfully integrated Groq Gemini LLM into a RAG pipeline with LangChain and FAISS.
 - Learned document chunking, embedding, and indexing for retrieval-based Q&A.
 - Handled multi-file uploads and user-friendly UI with Streamlit.
 - Solved PyTorch and embedding model loading issues related to device placement.
 - Built a job recommendation by skill matching from resumes and a job dataset.
 - Developed modular, maintainable Python code with clear separation of concerns.
 - Gained practical skills in a secure environment and API key management.
-